

# DATABASE SYSTEM PRINCIPLE - SQL QUERY LANGUAGE

李旭东  
LEEXUDONG@NANKAI.EDU.CN  
NANKAI UNIVERSITY

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## OVERVIEW OF THE SQL QUERY LANGUAGE

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory
  - Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
  - SQL-86, SQL-89, **SQL-92**
  - SQL:1999 (language name became Y2K compliant!)
  - SQL-2003, SQL-2006, SQL-2008
  - Not all examples here may work on your particular system.

©LXD

## OVERVIEW OF THE SQL QUERY LANGUAGE

- Structured Query Language (SQL)
  - Data-Definition Language, DDL
  - Data-Manipulation Language, DML
  - Integrity完整性
  - View Definition视图定义
  - Transaction Control事务控制
  - Embedded SQL and dynamic SQL嵌入式和动态SQL
  - Authorization授权
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- **Data Definition**
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## SQL DATA DEFINITION

- SQL DDL
  - The schema for each relation
  - The types of values associated with each attribute
  - The integrity constraints
  - The set of indices to be maintained for each relation
  - The security and authorization information for each relation
  - The physical storage structure of each relation on disk

©LXD

## DOMAIN TYPES IN SQL

- `char(n)`. Fixed length character string, with user-specified length  $n$ .
- `varchar(n)`. Variable length character strings, with user-specified maximum length  $n$ .
- `int`. Integer (a finite subset of the integers that is machine-dependent).
- `smallint`. Small integer (a machine-dependent subset of the integer domain type).
- `numeric(p,d)`. Fixed point number, with user-specified precision of  $p$  digits, with  $d$  digits to the right of decimal point. (ex., `numeric(3,1)`, allows 44.5, not 444.5 or 0.32)
- `real`, `double precision`. Floating point and double-precision floating point numbers, with machine-dependent precision.
- `float(n)`. Floating point number, with user-specified precision of at least  $n$  digits.

©LXD

## BASIC SCHEMA DEFINITION

- Create table
- Insert into
- Delete from
- Alter table
- Drop table

©LXD

## CREATE TABLE CONSTRUCT

- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1,
               A2 D2, ...,
               An Dn,
               (integrity-constraint1 完整性约束),
               ...,
               (integrity-constraintk))
```

- $r$  is the name of the relation
- each  $A_i$  is an attribute name in the schema of relation  $r$
- $D_i$  is the data type of values in the domain of attribute  $A_i$

©LXD

## CREATE TABLE CONSTRUCT

- the **create table** command
- Example:

```
create table instructor (
    ID      char(5),
    name    varchar(20),
    dept_name varchar(20),
    salary  numeric(8,2));
```

©LXD

## INTEGRITY CONSTRAINTS IN CREATE TABLE

- **not null**
- **primary key** ( $A_1, \dots, A_n$ )
- **foreign key** ( $A_m, \dots, A_n$ ) **references**  $r$

Example:

```
create table instructor (
    ID      char(5),
    name    varchar(20) not null,
    dept_name varchar(20),
    salary  numeric(8,2),
    primary key (ID),
    foreign key (dept_name) references department);
```

**primary key** declaration on an attribute automatically ensures **not null**

©LXD

## CASES: CREATE TABLE

```
create table department
(dept_name varchar (20),
building varchar (15),
budget numeric (12,2),
primary key (dept name);
```

```
create table course
(course id varchar (7),
title varchar (50),
credits numeric (2,0),
primary key (course id),
foreign key (dept name) references department);
```

©LXD

## CASES: CREATE TABLE

```
create table section
(course id varchar (8),
sec id varchar (8),
semester varchar (6),
year numeric (4,0),
building varchar (15),
room number varchar (7),
time slot id varchar (4),
primary key (course id, sec id, semester, year),
foreign key (course id) references course);
```

©LXD

## CASES: CREATE TABLE

```
create table teaches
( ID varchar (5),
course id varchar (8),
sec id varchar (8),
semester varchar (6),
year numeric (4,0),
primary key ( ID , course id, sec id, semester, year),
foreign key (course id, sec id, semester, year) references section,
foreign key ( ID ) references instructor);
```

©LXD

## UPDATES TO TABLES

- **Insert**
  - **insert into** *instructor* values ('10211', 'Smith', 'Biology', 66000);
- **Delete**
  - Remove all tuples from the *student* relation
    - **delete from** *student*
- **Drop Table**
  - **drop table** *r*

©LXD

## UPDATES TO TABLES

- **Alter**
  - **alter table** *r* **add** *A D*
    - where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*.
    - All exiting tuples in the relation are assigned *null* as the value for the new attribute.
  - **alter table** *r* **drop** *A*
    - where *A* is the name of an attribute of relation *r*
    - Dropping of attributes not supported by many databases.

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- **Basic Query Structure**
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## BASIC STRUCTURE OF SQL QUERIES

- A typical SQL query has the form:

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

*A*<sub>*i*</sub> represents an attribute

- *R*<sub>*i*</sub> represents a relation
- *P* is a predicate.

- The result of an SQL query is a **relation**.

©LXD



## QUERIES ON A SINGLE RELATION

- The **select** clause lists the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra
- Example: find the names of all instructors:
 

```
select name
from instructor
```
- NOTE: SQL names are case insensitive
  - E.g., *Name*  $\equiv$  *NAME*  $\equiv$  *name*
  - Some people use upper case wherever we use bold font.

name
Srinivasan
Wu
Mozart
Einstein
El Said
Gold
Katz
Califieri
Singh
Crick
Brandt
Kim

©LXD

## QUERIES ON A SINGLE RELATION

### - DEDUPLICATION 去重

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after select.
- Find the department names of all instructors, and remove duplicates

```
select distinct dept_name
from instructor
```

- The keyword **all** specifies that duplicates should not be removed.
 

```
select all dept_name
from instructor
```

©LXD

## QUERIES ON A SINGLE RELATION

### - MORE

- An **asterisk** in the select clause denotes "all attributes"
 

```
select *
from instructor
```
- An attribute can be a **literal** with no **from** clause
 

```
select '437'
```

  - Results is a table with one column and a single row with value "437"
  - Can give the column a name using:
 

```
select '437' as FOO
```

©LXD

## QUERIES ON A SINGLE RELATION

### - MORE

- An attribute can be a literal with **from** clause
 

```
select 'A'
from instructor
```

  - Result is a table with **one column** and **N rows** (number of tuples in the *instructors* table), each row with value "A"

©LXD

## QUERIES ON A SINGLE RELATION

### - MORE

- The **select** clause can contain arithmetic expressions involving the operation, +, -, \*, and /, and operating on constants or attributes of tuples.
  - The query:
 

```
select ID, name, salary/12
from instructor
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.
  - Can rename "*salary/12*" using the **as** clause:
 

```
select ID, name, salary/12 as monthly_salary
```

©LXD

## THE WHERE CLAUSE

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.
- To find all instructors in Comp. Sci. dept

```
select name
from instructor
where dept_name = 'Comp. Sci.'
```

©LXD

## THE WHERE CLAUSE

- Comparison results can be combined using the logical connectives **and**, **or**, and **not**
  - To find all instructors in Comp. Sci. dept with salary > 80000
 

```
select name
from instructor
where dept_name = 'Comp. Sci.' and salary > 80000
```
- Comparisons can be applied to results of arithmetic expressions.

©LXD

## QUERIES ON MULTIPLE RELATIONS

```
select A1 , A2 ,..., An
from r1 , r2 ,...,rm
where P;
```

- Each  $A_i$  represents an attribute, and each  $r_i$  a relation.  $P$  is a predicate.
  - If the where clause is omitted, the predicate  $P$  is true.

©LXD

## QUERIES ON MULTIPLE RELATIONS

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product *instructor*  $\times$  *teaches*

```
select *
from instructor, teaches
```

  - generates every possible instructor – teaches pair, with all attributes from both relations.
  - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)

©LXD

## CARTESIAN PRODUCT 笛卡尔积

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

inst_ID	name	dept_name	salary	teaches_ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...

©LXD

## CASES: CARTESIAN PRODUCT

- Find the names of all instructors who have taught some course and the course\_id
  - ```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID
```
- Find the names of all instructors in the Art department who have taught some course and the course\_id
  - ```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID and instructor.dept_name = 'Art'
```

©LXD

## THE NATURAL JOIN

```
select A1, A2, ..., An
from r1 natural join r2 natural join ... natural join r_m
where P;
```

More generally, a **from** clause can be of the form

**from**  $E_1, E_2, \dots, E_n$

- where each  $E_i$  can be a single relation or an expression involving natural joins.

©LXD

## CASES: THE NATURAL JOIN

The natural join of the instructor relation with the teaches relation

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

ID	name	dept_name	salary	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
12121	Wu	Finance	90000	FIN-201	1	Spring	2010
15151	Mozart	Music	40000	MU-199	1	Spring	2010
22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
32343	El Said	History	60000	HIS-351	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
76766	Crick	Biology	72000	BIO-101	1	Summer	2009
76766	Crick	Biology	72000	BIO-301	1	Summer	2010
83821	Brandt	Comp. Sci.	92000	CS-190	1	Spring	2009
83821	Brandt	Comp. Sci.	92000	CS-190	2	Spring	2009
83821	Brandt	Comp. Sci.	92000	CS-319	2	Spring	2010
98345	Kim	Elec. Eng.	80000	EE-181	1	Spring	2009

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

©LXD

## CASES: THE NATURAL JOIN

## • Way 1

```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID;
```

## • Way 2

```
select name, course_id
from instructor natural join teaches;
```

©LXD

## CASES 2: THE NATURAL JOIN

- ```
select name, title
from instructor natural join teaches, course
where teaches.course_id = course.course_id;
```
- ```
select name, title
from instructor natural join teaches natural join course;
```
- ```
select name, title
from (instructor natural join teaches) join course using (course_id);
```

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- **Additional Basic Operations**
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## THE RENAME OPERATION

- The SQL allows renaming relations and attributes using the **as** clause:  
*old-name as new-name*
- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.
  - ```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Comp. Sci.'
```
- Keyword **as** is optional and may be omitted  
*instructor as T  $\equiv$  instructor T*

©LXD

## STRING OPERATIONS

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring.
  - underscore ( \_ ). The \_ character matches any character.

©LXD



## STRING OPERATIONS

- Find the names of all instructors whose name includes the substring "dar".

```
select name
from instructor
where name like '%dar%'
```

- Match the string "100%"

```
like '100 \%' escape '\'
```

in that above we use backslash (\) as the escape character.

©LXD

## STRING OPERATIONS

- Patterns are case sensitive.
- Pattern matching examples:
  - 'Intro%' matches any string beginning with "Intro".
  - '%Comp%' matches any string containing "Comp" as a substring.
  - '\_\_\_' matches any string of exactly three characters.
  - '\_\_\_%' matches any string of at least three characters.

©LXD

## STRING OPERATIONS

- SQL supports a variety of string operations such as
  - concatenation (using "||")
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.

©LXD

## ORDERING THE DISPLAY OF TUPLES

- List in alphabetic order the names of all instructors
 

```
select distinct name
from instructor
order by name
```
- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
  - Example: **order by name desc**
- Can sort on multiple attributes
  - Example: **order by dept\_name, name**

©LXD

## WHERE CLAUSE PREDICATES

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq 90,000$  and  $\leq 100,000$ )
 

```
select name
from instructor
where salary between 90000 and 100000
```

©LXD

## WHERE CLAUSE PREDICATES

- Way 1:
 

```
select name, course_id
from instructor, teaches
where (instructor.ID= teaches.ID and dept_name = 'Biology');
```
- Way 2:
 

```
select name, course_id
from instructor, teaches
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- **Set Operations**
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## SET OPERATIONS

- union
- intersect
- except

©LXD

## SET OPERATIONS

## - UNION

- Find courses that ran in Fall 2009 or in Spring 2010

```
(select course_id from section where sem = 'Fall' and
year = 2009)
union
(select course_id from section where sem = 'Spring' and
year = 2010)
```

©LXD

## SET OPERATIONS

## - INTERSECT

- Find courses that ran in Fall 2009 and in Spring 2010

```
(select course_id from section where sem = 'Fall' and
year = 2009)
intersect
(select course_id from section where sem = 'Spring'
and year = 2010)
```

©LXD

## SET OPERATIONS

## - EXCEPT

- Find courses that ran in Fall 2009 but not in Spring 2010

```
(select course_id from section where sem = 'Fall' and
year = 2009)
except
(select course_id from section where sem = 'Spring'
and year = 2010)
```

©LXD

## SET OPERATIONS

- Set operations **union**, **intersect**, and **except**

- Each of the above operations automatically eliminates duplicates

- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.

- Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:

- $m + n$  times in  $r$  **union all**  $s$
- $\min(m, n)$  times in  $r$  **intersect all**  $s$
- $\max(0, m - n)$  times in  $r$  **except all**  $s$

©LXD



## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- **Null Values**
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## NULL VALUES

- It is possible for tuples to have a null value, denoted by **null**; for some of their attributes
- **null** signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving **null** is **null**
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name
from instructor
where salary is null
```

©LXD

## NULL VALUES AND THREE VALUED LOGIC

- Three values – **true, false, unknown**
- Any comparison with **null** returns **unknown**
  - Example:  $5 < \text{null}$  or  $\text{null} <> \text{null}$  or  $\text{null} = \text{null}$

©LXD

## NULL VALUES AND THREE VALUED LOGIC

- Three-valued logic using the value **unknown**:
  - OR:  $(\text{unknown or true}) = \text{true}$ ,  
 $(\text{unknown or false}) = \text{unknown}$   
 $(\text{unknown or unknown}) = \text{unknown}$
  - AND:  $(\text{true and unknown}) = \text{unknown}$ ,  
 $(\text{false and unknown}) = \text{false}$ ,  
 $(\text{unknown and unknown}) = \text{unknown}$
  - NOT:  $(\text{not unknown}) = \text{unknown}$
  - “**P is unknown**” evaluates to true if predicate *P* evaluates to **unknown**
- Result of **where** clause predicate is treated as **false** if it evaluates to **unknown**

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- **Aggregate Functions**
- Nested Subqueries
- Modification of the Database

©LXD

## AGGREGATE FUNCTIONS 聚集函数

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg**: average value  
**min**: minimum value  
**max**: maximum value  
**sum**: sum of values  
**count**: number of values

©LXD

## AGGREGATE FUNCTIONS

- Find the average salary of instructors in the Computer Science department
  - `select avg (salary)`  
`from instructor`  
`where dept_name= 'Comp. Sci.';`
- Find the total number of instructors who teach a course in the Spring 2010 semester
  - `select count (distinct ID)`  
`from teaches`  
`where semester = 'Spring' and year = 2010;`

©LXD

## AGGREGATE FUNCTIONS – GROUP BY

- Find the average salary of instructors in each department
  - `select dept_name, avg (salary) as avg_salary`  
`from instructor`  
`group by dept_name;`

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76943	Singh	Finance	80000
32343	El Said	History	60000
88583	Califrieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

dept_name	salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000

©LXD

## AGGREGATE FUNCTIONS – GROUP BY

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list
  - `/* erroneous query */`  
`select dept_name, ID, avg (salary)`  
`from instructor`  
`group by dept_name;`

©LXD

## AGGREGATE FUNCTIONS – HAVING CLAUSE

- Find the names and average salaries of all departments whose average salary is greater than 42000
  - `select dept_name, avg (salary)`  
`from instructor`  
`group by dept_name`  
`having avg (salary) > 42000;`

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

©LXD

## AGGREGATE FUNCTIONS

- Total all salaries
  - `select sum (salary )`  
`from instructor`
  - Above statement ignores null amounts
  - Result is *null* if there is no non-null amount

©LXD

## AGGREGATE FUNCTIONS

- All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
  - count returns 0
  - all other aggregates return null

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- **Nested Subqueries**
- Modification of the Database

©LXD

## NESTED SUBQUERIES 嵌套子查询

- SQL provides a mechanism for nesting subqueries.
- A subquery is a select-from-where expression that is nested within another query.

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
course_id in (select course_id
from section
where semester = 'Spring' and year= 2010);
```

©LXD

## NESTED SUBQUERIES 嵌套子查询

- The nesting can be done in the following SQL query
- ```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

as follows:

- A<sub>i</sub> can be replaced by a subquery that generates a single value.
- r<sub>i</sub> can be replaced by any valid subquery
- P can be replaced with an expression of the form:

$$B \langle \text{operation} \rangle (\text{subquery})$$

Where B is an attribute and &lt;operation&gt; to be defined later.

©LXD

## SUBQUERIES IN THE WHERE CLAUSE

- A common use of subqueries is to perform tests:
  - For set membership
  - For set comparisons
  - For set cardinality.

©LXD

## NESTED SUBQUERIES

## - SET MEMBERSHIP 集合成员测试

- in, not in

- Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
course_id not in (select course_id
from section
where semester = 'Spring' and year= 2010);
```

©LXD

## NESTED SUBQUERIES

## - SET MEMBERSHIP

- Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101

```
select count (distinct ID)
from takes
where (course_id, sec_id, semester, year) in
(select course_id, sec_id, semester, year
from teaches
where teaches.ID= 10101);
```

©LXD



## NESTED SUBQUERIES

## - SET COMPARISON 集合比较

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Biology';
```

- Same query using **> some** clause

```
select name
from instructor
where salary > some (select salary
from instructor
where dept_name = 'Biology');
```

©LXD

## NESTED SUBQUERIES

## - SET COMPARISON 集合比较

- > some**
- < some**
- <= some**
- >= some**
- <> some**
- = some** (same to **in**)

•  $F < \text{comp} > \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F < \text{comp} > t)$   
Where  $< \text{comp} >$  can be:  $<, \leq, >, \geq, \neq$

|           |                                                                                       |   |   |                      |        |                                        |
|-----------|---------------------------------------------------------------------------------------|---|---|----------------------|--------|----------------------------------------|
| (5 < some | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr><tr><td>6</td></tr></table> ) | 0 | 5 | 6                    | = true | (read: 5 < some tuple in the relation) |
| 0         |                                                                                       |   |   |                      |        |                                        |
| 5         |                                                                                       |   |   |                      |        |                                        |
| 6         |                                                                                       |   |   |                      |        |                                        |
| (5 < some | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr></table> )                    | 0 | 5 | = false              |        |                                        |
| 0         |                                                                                       |   |   |                      |        |                                        |
| 5         |                                                                                       |   |   |                      |        |                                        |
| (5 = some | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr></table> )                    | 0 | 5 | = true               |        |                                        |
| 0         |                                                                                       |   |   |                      |        |                                        |
| 5         |                                                                                       |   |   |                      |        |                                        |
| (5 ≠ some | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr></table> )                    | 0 | 5 | = true (since 0 ≠ 5) |        |                                        |
| 0         |                                                                                       |   |   |                      |        |                                        |
| 5         |                                                                                       |   |   |                      |        |                                        |

(= some) = in  
However, (≠ some) ≠ not in

©LXD

## NESTED SUBQUERIES

## - SET COMPARISON 集合比较

- > all**
- < all**
- <= all**
- >= all**
- = all** (not same to **in**)
- <> all** (same to **not in**)

•  $F < \text{comp} > \text{all } r \Leftrightarrow \forall t \in r (F < \text{comp} > t)$

|          |                                                                                       |   |    |                                |         |  |
|----------|---------------------------------------------------------------------------------------|---|----|--------------------------------|---------|--|
| (5 < all | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr><tr><td>6</td></tr></table> ) | 0 | 5  | 6                              | = false |  |
| 0        |                                                                                       |   |    |                                |         |  |
| 5        |                                                                                       |   |    |                                |         |  |
| 6        |                                                                                       |   |    |                                |         |  |
| (5 < all | <table border="1"><tr><td>6</td></tr><tr><td>10</td></tr></table> )                   | 6 | 10 | = true                         |         |  |
| 6        |                                                                                       |   |    |                                |         |  |
| 10       |                                                                                       |   |    |                                |         |  |
| (5 = all | <table border="1"><tr><td>4</td></tr><tr><td>5</td></tr></table> )                    | 4 | 5  | = false                        |         |  |
| 4        |                                                                                       |   |    |                                |         |  |
| 5        |                                                                                       |   |    |                                |         |  |
| (5 ≠ all | <table border="1"><tr><td>4</td></tr><tr><td>6</td></tr></table> )                    | 4 | 6  | = true (since 5 ≠ 4 and 5 ≠ 6) |         |  |
| 4        |                                                                                       |   |    |                                |         |  |
| 6        |                                                                                       |   |    |                                |         |  |

(= all) = not in  
However, (≠ all) ≠ in

©LXD

## NESTED SUBQUERIES

## - SET COMPARISON 集合比较

?What meaning?

```
select dept_name
from instructor
group by dept_name
having avg (salary) >= all (select avg (salary)
from instructor
group by dept_name);
```

©LXD

## TEST FOR EMPTY RELATIONS

- exists , not exists**
  - Yet another way of specifying the query "Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester"
- ```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
exists (select *
from section as T
where semester = 'Spring' and year = 2010 and S.course_id = T.course_id);
```
- Correlation name** – variable S in the outer query
  - Correlated subquery** – the inner query

©LXD

## TEST FOR EMPTY RELATIONS

- Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name
from student as S
where not exists ((select course_id
from course
where dept_name = 'Biology')
except
(select T.course_id
from takes as T
where S.ID = T.ID));
```

- First nested query lists all courses offered in Biology
- Second nested query lists all courses a particular student took

■ Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$

■ Note: Cannot write this query using = all and its variants

©LXD

### TEST FOR ABSENCE OF DUPLICATE TUPLES

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- The **unique** construct evaluates to "true" if a given subquery contains no duplicates.
- Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique (select R.course_id
              from section as R
              where T.course_id = R.course_id
              and R.year = 2009);
```

©LXD

### SUBQUERIES IN THE FROM CLAUSE

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
      from instructor
      group by dept_name)
where avg_salary > 42000;
```

- Note that we do not need to use the **having** clause

©LXD

### SUBQUERIES IN THE FROM CLAUSE

- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."
- Another way to write above query

```
select dept_name, avg_salary
from (select dept_name, avg (salary)
      from instructor
      group by dept_name) as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```

©LXD

### WITH CLAUSE

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs
- Introduced in SQL:1999
- Find all departments with the maximum budget

```
with max_budget (value) as
(select max(budget)
 from department)
select department.name
from department, max_budget
where department.budget = max_budget.value;
```

©LXD

### COMPLEX QUERIES USING WITH CLAUSE

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as
(select dept_name, sum(salary)
 from instructor
 group by dept_name),
dept_total_avg(value) as
(select avg(value)
 from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value > dept_total_avg.value;
```

©LXD

### SCALAR SUBQUERIES 标量子查询

- Scalar subqueries
  - SQL allows subqueries to occur wherever an expression returning a value is permitted, provided the subquery returns only one tuple containing a single attribute
  - List all departments along with the number of instructors in each department

```
select dept_name,
(select count(*)
 from instructor where department.dept_name = instructor.dept_name)
as num_instructors
from department;
```

- Runtime error if subquery returns more than one result tuple

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- **Modification of the Database**

©LXD

## MODIFICATION OF THE DATABASE

- Deletion of tuples from a given relation.
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation

©LXD

## DELETION 删除

- Delete all instructors
 

```
delete from instructor
```
- Delete all instructors from the Finance department
 

```
delete from instructor
where dept_name = 'Finance';
```
- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.
 

```
delete from instructor
where dept_name in (select dept_name
from department
where building = 'Watson');
```

©LXD

## DELETION

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor
where salary < (select avg (salary)
from instructor);
```

- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
  1. First, compute **avg** (salary) and find all tuples to delete
  2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

©LXD

## INSERTION 插入

- Add a new tuple to *course*

```
insert into course
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```
- or equivalently
 

```
insert into course (course_id, title, dept_name, credits)
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```
- Add a new tuple to *student* with *tot\_creds* set to null
 

```
insert into student
values ('3003', 'Green', 'Finance', null);
```

©LXD

## INSERTION

- Add all instructors to the *student* relation with *tot\_creds* set to 0

```
insert into student
select ID, name, dept_name, 0
from instructor
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

Otherwise queries like

```
insert into table1 select * from table1
```

would cause problem

©LXD



## UPDATES更新

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
  - Write two **update** statements:
 

```
update instructor
set salary = salary * 1.03
where salary > 100000;
update instructor
set salary = salary * 1.05
where salary <= 100000;
```
  - *The order is important*
  - Can be done better using the **case** statement (next slide)

©LXD

## UPDATE

### - CASE STATEMENT FOR CONDITIONAL UPDATES

- Same query as before but with **case** statement
 

```
update instructor
set salary = case
when salary <= 100000 then salary * 1.05
else salary * 1.03
end
```

©LXD

## UPDATE

### - CASE STATEMENT FOR CONDITIONAL UPDATES

- The general form of the case statement is as follows

```
case
when pred1 then result1
when pred2 then result2
...
when predn then resultn
else result0
end
```

©LXD

## UPDATES WITH SCALAR SUBQUERIES

- Recompute and update **tot\_creds** value for all students
 

```
update student S
set tot_cred = (select sum(credits)
from takes, course
where takes.course_id = course.course_id and
S.ID= takes.ID and takes.grade <> 'F' and
takes.grade is not null);
```
- Sets **tot\_creds** to null for students who have not taken any course

©LXD

## UPDATES WITH SCALAR SUBQUERIES

- Recompute and update **tot\_creds** value for all students
- Instead of **sum(credits)**, use:

```
update student S
set tot_cred =
case
when sum(credits) is not null then sum(credits)
else 0
end
```

©LXD

## SUMMARY

©LXD

Q&A?

THANKS!

leexudong@nankai.edu.cn