



DATABASE SYSTEM PRINCIPLE — ADVANCED SQL

李旭东

LEEXUDONG@NANKAI.EDU.CN

NANKAI UNIVERSITY

OBJECTIVES

- Accessing SQL from a Programming Language
- Temporary Table
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP

OBJECTIVES

- Accessing SQL from a Programming Language
- Temporary Table
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP

ACCESSING SQL FROM A PROGRAMMING LANGUAGE

- API (application-program interface) for a program to interact with a database server
- Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables

ACCESSING SQL FROM A PROGRAMMING LANGUAGE

- Various tools:
 - JDBC (Java Database Connectivity) works with Java
 - ODBC (Open Database Connectivity) works with C, C++, C#, ...
 - Embedded SQL

JDBC

- **JDBC** is a Java API for communicating with database systems supporting SQL.
- JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes.

JDBC

- Model for communicating with the database:
 - Open a connection
 - Create a “statement” object
 - Execute queries using the Statement object to send queries and fetch results
 - Exception mechanism to handle errors

JDBC: JAVA VM + ECLIPSE IDE + MYSQL DB

- JAVA VM(Java Platform (JDK) 8u121):

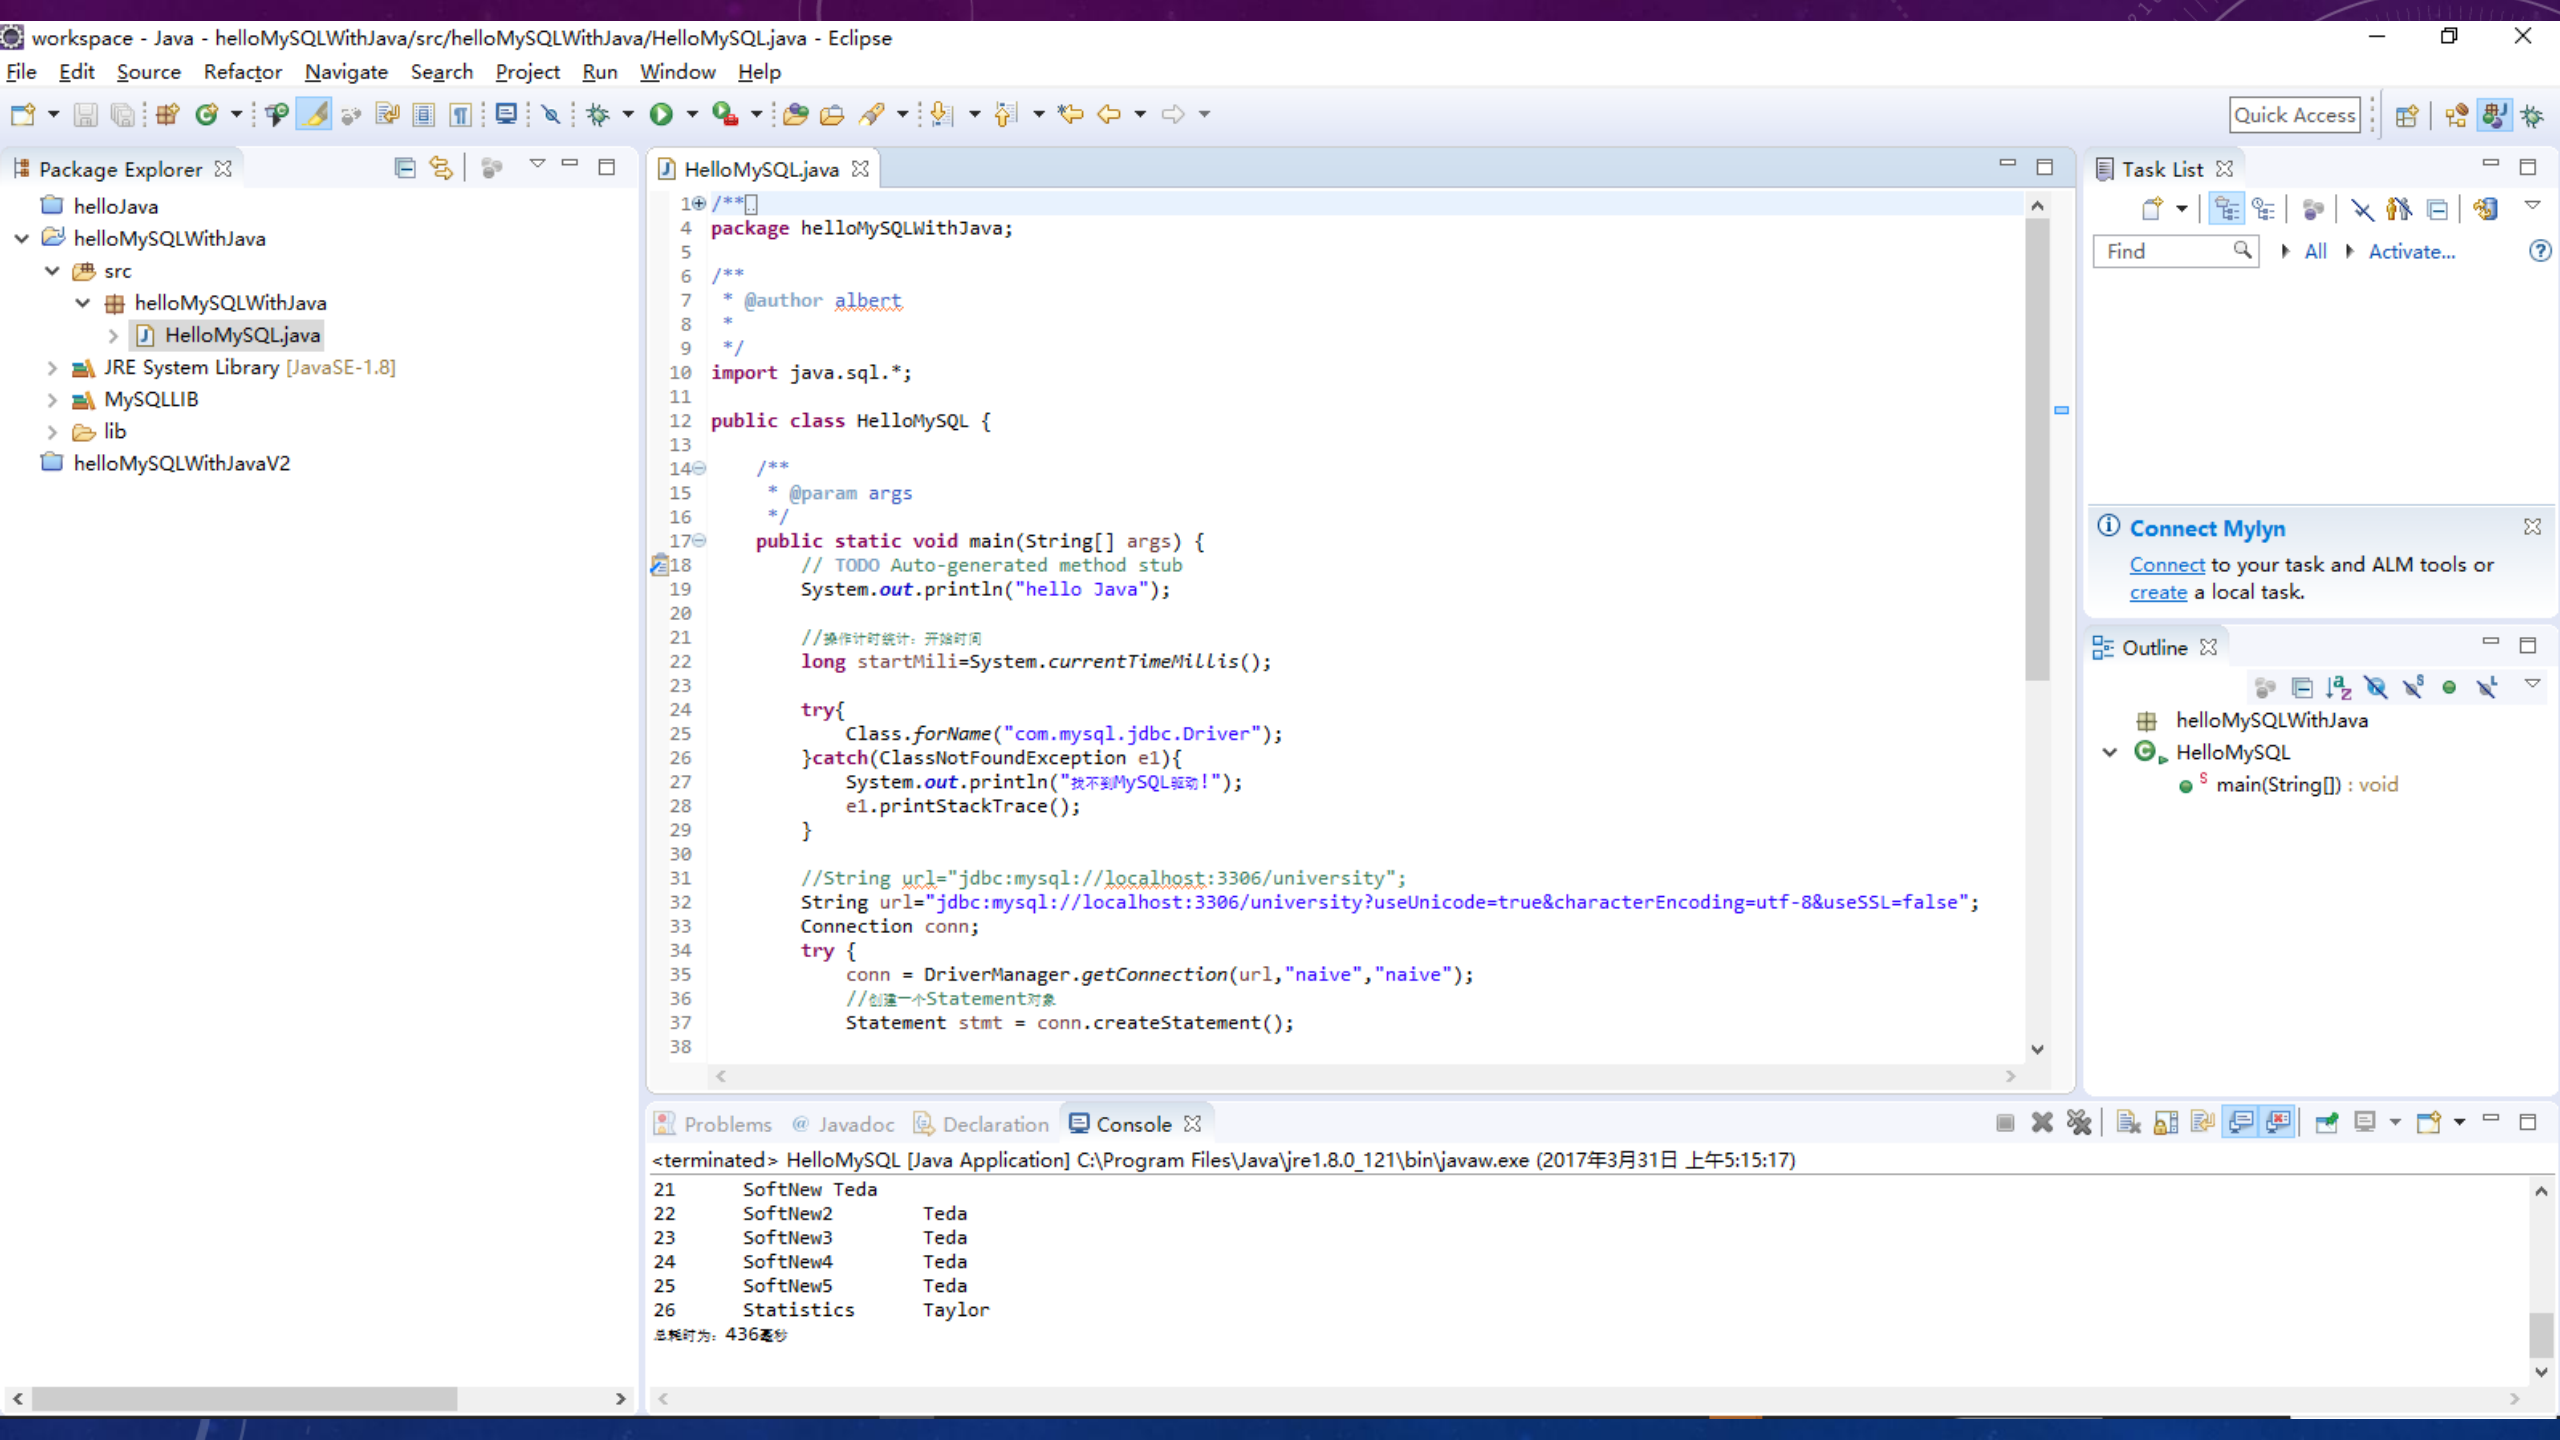
<http://www.oracle.com/technetwork/java/javase/downloads/index-jsp-138363.html>

- eclipse IDE(Eclipse Neon):

<http://www.eclipse.org/downloads/>

- MySQL Connector for Java (mysql-connector-java-5.1.41.zip):

<https://dev.mysql.com/downloads/connector/j/>



JDBC

```
public static void JDBCexample(String userid, String passwd)
{
    try
    {
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@db.yale.edu:1521:univdb",
            userid, passwd);
        Statement stmt = conn.createStatement();
        try {
            stmt.executeUpdate(
                "insert into instructor values('77987', 'Kim', 'Physics', 98000)");
        } catch (SQLException sqle)
        {
            System.out.println("Could not insert tuple. " + sqle);
        }
        ResultSet rset = stmt.executeQuery(
            "select dept_name, avg (salary) "+
            " from instructor "+
            " group by dept_name");
        while (rset.next()) {
            System.out.println(rset.getString("dept_name") + " " +
                rset.getFloat(2));
        }
        stmt.close();
        conn.close();
    }
    catch (Exception sqle)
    {
        System.out.println("Exception : " + sqle);
    }
}
```

JDBC

```
"insert into instructor values(' " + ID + " ', ' " + name + " ', " +  
" ' + dept_name + " ', " ' balance + ")"
```


JDBC: SQL INJECTION

```
"select * from instructor where name = '" + name + "'"
```

If the user, instead of entering a name, enters:

X' or 'Y' = 'Y

then the resulting statement becomes:

```
"select * from instructor where name = '" + "X' or 'Y' = 'Y'" + "'"
```

which is:

```
select * from instructor where name = 'X' or 'Y' = 'Y'
```

JDBC: PREPARED STATEMENTS

```
PreparedStatement pstmt = conn.prepareStatement(  
    "insert into instructor values(?,?,?,?)");  
pstmt.setString(1, "88877");  
pstmt.setString(2, "Perry");  
pstmt.setString(3, "Finance");  
pstmt.setInt(4, 125000);  
pstmt.executeUpdate();  
pstmt.setString(1, "88878");  
pstmt.executeUpdate();
```

No args

JDBC: PREPARED STATEMENTS

- Blob, Clob
- PreparedStatement
 - `setBlob(int parameterIndex, InputStream inputStream)`
 - `getBlob()`

JDBC: CALLABLE STATEMENTS

可调用语句

```
String sql = "call sp_department1(?,?);";
```

```
CallableStatement cs = null;
```

```
cs = conn.prepareCall(sql);
```

```
cs.setString(1, "Soft");
```

```
cs.registerOutParameter(2, Types.INTEGER);
```

```
cs.execute();
```

```
int mCnt = cs.getInt(2);
```

```
cs.close();
```

JDBC: METADATA FEATURES

```
ResultSetMetaData rsmd = rs.getMetaData();  
for(int i = 1; i <= rsmd.getColumnCount(); i++) {  
    System.out.println(rsmd.getColumnName(i));  
    System.out.println(rsmd.getColumnTypeName(i));  
}
```

JDBC: METADATA FEATURES

```
DatabaseMetaData dbmd = conn.getMetaData();
ResultSet rs = dbmd.getColumns(null, "univdb", "department", "%");
    // Arguments to getColumns: Catalog, Schema-pattern, Table-pattern,
    //                          and Column-Pattern
    // Returns: One row for each column; row has a number of attributes
    //          such as COLUMN_NAME, TYPE_NAME
while( rs.next()) {
    System.out.println(rs.getString("COLUMN_NAME"),
                       rs.getString("TYPE_NAME");
}
```


ODBC

- Open DataBase Connectivity (ODBC) standard
 - standard for application program to communicate with a database server.
 - application program interface (API) to
 - open a connection with a database,
 - send queries and updates,
 - get back results.
- Applications such as GUI, spreadsheets, etc. can use ODBC

ODBC

```
void ODBCexample()
{
    RETCODE error;
    HENV env; /* environment */
    HDBC conn; /* database connection */

    SQLAllocEnv(&env);
    SQLAllocConnect(env, &conn);
    SQLConnect(conn, "db.yale.edu", SQL_NTS, "avi", SQL_NTS,
               "avipasswd", SQL_NTS);
    {
        char deptname[80];
        float salary;
        int lenOut1, lenOut2;
        HSTMT stmt;

        char * sqlquery = "select dept_name, sum (salary)
                           from instructor
                           group by dept_name";
        SQLAllocStmt(conn, &stmt);
        error = SQLExecDirect(stmt, sqlquery, SQL_NTS);
        if (error == SQL_SUCCESS) {
            SQLBindCol(stmt, 1, SQL_C_CHAR, deptname, 80, &lenOut1);
            SQLBindCol(stmt, 2, SQL_C_FLOAT, &salary, 0, &lenOut2);
            while (SQLFetch(stmt) == SQL_SUCCESS) {
                printf (" %s %g\n", deptname, salary);
            }
        }
        SQLFreeStmt(stmt, SQL_DROP);
    }
    SQLDisconnect(conn);
    SQLFreeConnect(conn);
    SQLFreeEnv(env);
}
```

EMBEDDED SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, C++, Java, Fortran, and PL/1,
- A language to which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise *embedded SQL*.
- The basic form of these languages follows that of the System R embedding of SQL into PL/1.

EMBEDDED SQL

- **EXEC SQL** statement is used to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement >;

Note: this varies by language:

- In some languages, like COBOL, the semicolon is replaced with END-EXEC
- In Java embedding uses # SQL { };

EMBEDDED SQL

- Before executing any SQL statements, the program must first connect to the database. This is done using:

EXEC-SQL connect to *server* user *user-name* using *password*;

Here, *server* identifies the server to which a connection is to be established.

- Variables of the host language can be used within embedded SQL statements. They are preceded by a colon (:) to distinguish from SQL variables (e.g., *:credit_amount*)

EMBEDDED SQL

- Variables used as above must be declared within DECLARE section, as illustrated below. The syntax for declaring the variables, however, follows the usual host language syntax.

```
EXEC-SQL BEGIN DECLARE SECTION}
```

```
    int credit-amount ;
```

```
EXEC-SQL END DECLARE SECTION;
```


EMBEDDED SQL

- To write an embedded SQL query, we use the **declare *c* cursor for <SQL query>** statement. The variable *c* is used to identify the query
- Example:
 - From within a host language, find the ID and name of students who have completed more than the number of credits stored in variable **credit_amount** in the host language

EMBEDDED SQL

- Example: (cont.,)
 - Specify the query in SQL as follows:

EXEC SQL

declare *c* cursor for

select *ID, name*

from *student*

where *tot_cred > :credit_amount*

END_EXEC

EMBEDDED SQL

- The **open** statement for our example is as follows:

```
EXEC SQL open c ;
```

This statement causes the database system to execute the query and to save the results within a temporary relation. The query uses the value of the host-language variable *credit-amount* at the time the **open** statement is executed.

- The fetch statement causes the values of one tuple in the query result to be placed on host language variables.

```
EXEC SQL fetch c into :si, :sn END_EXEC
```

Repeated calls to fetch get successive tuples in the query result

EMBEDDED SQL

- A variable called SQLSTATE in the SQL communication area (SQLCA) gets set to '02000' to indicate no more data is available
- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

EXEC SQL close c ;

Note: above details vary with language. For example, the Java embedding defines Java iterators to step through result tuples.

UPDATES THROUGH EMBEDDED SQL

- Embedded SQL expressions for database modification (**update**, **insert**, and **delete**)
- Can update tuples fetched by cursor by declaring that the cursor is for update

EXEC SQL

```
declare c cursor for  
select *  
from instructor  
where dept_name = 'Music'  
for update
```


UPDATES THROUGH EMBEDDED SQL

- Embedded SQL expressions for database modification (**update**, **insert**, and **delete**)
- We then iterate through the tuples by performing **fetch** operations on the cursor (as illustrated earlier), and after fetching each tuple we execute the following code:

```
update instructor  
set salary = salary + 1000  
where current of c
```

SQLJ

```
#sql iterator deptInfolter ( String dept_name, int avgSal);
deptInfolter iter = null;

#sql iter = { select dept_name, avg(salary)
              from instructor
              group by dept_name };
while (iter.next()) {
    String deptName = iter.dept_name();
    int avgSal = iter.avgSal();
    System.out.println(deptName + " " + avgSal);
}
iter.close();
```

PYTHON

```
import pymysql

conn=pymysql.connect(host='localhost',user='user',password='pwd',database='schema1',charset='utf8')

cursor=conn.cursor()

sql = "select name,salary,dept_name from instructor where dept_name like '%s'" %("C"+"%")

cursor.execute(sql)

retdat = cursor.fetchall()

for row in retdat:

    print(row[0], ': ', row[1], ': ', row[2])

cursor.close()

conn.close()
```

*Pymysql error: cryptography is required for sha256_password
or caching_sha2_password
pip install **cryptography***

OBJECTIVES

- Accessing SQL from a Programming Language
- **Temporary Table**
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP

TEMPORARY TABLE

- Temporary Table
 - a temporary table is a special type of table that allows you to store a temporary result set, which you can reuse several times in a single session
 - MySQL removes the temporary table automatically when the session ends or the connection is terminated
 - A temporary table can have the same name as a normal table in a database. For example, if you create a temporary table named employees in the sample database, the existing employees table becomes inaccessible.

TEMPORARY TABLE

- Temporary Table
 - A temporary table is only available and accessible to the client that creates it.
 - Different clients can create temporary tables with the same name without causing errors because only the client that creates the temporary table can see it.
 - However, in the same session, two temporary tables cannot share the same name.

TEMPORARY TABLE – CREATE

```
CREATE TEMPORARY TABLE top10customers
```

```
SELECT p.customerNumber, c.customerName,
```

```
       ROUND(SUM(p.amount), 2) sales
```

```
FROM payments p
```

```
INNER JOIN customers c ON c.customerNumber = p.customerNumber
```

```
GROUP BY p.customerNumber
```

```
ORDER BY sales DESC
```

```
LIMIT 10;
```

round(x,d) :四舍五入
x指要处理的数,
d是指保留几位小数

TEMPORARY TABLE – RETRIEVE

SELECT

customerNumber,
customerName,
sales

FROM

top10customers

ORDER BY sales;

TEMPORARY TABLE – DROP

```
DROP TEMPORARY TABLE top10customers;
```

OBJECTIVES

- Accessing SQL from a Programming Language
- Temporary Table
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP

FUNCTIONS AND PROCEDURES

- SQL:1999 supports functions and procedures
 - Functions/procedures can be written in SQL itself, or in an external programming language
 - Functions written in an external languages are particularly useful with specialized data types such as images and geometric objects.
 - Example: functions to check if polygons overlap, or to compare images for similarity.
 - **Some** database systems support **table-valued functions**, which can return a relation as a result.

FUNCTIONS AND PROCEDURES

- SQL:1999 also supports a rich set of imperative constructs, including
 - Loops, if-then-else, assignment
- Many databases have proprietary procedural extensions to SQL that differ from SQL:1999.
- PL/SQL (oracle), TransactSQL (ms sql server), PL/pgSQL

SQL FUNCTIONS

- Define a function that, given the name of a department, returns the count of the number of instructors in that department.

```
create function dept_count (dept_name varchar(20))  
  returns integer  
  begin  
    declare d_count integer;  
    select count ( * ) into d_count  
    from instructor  
    where instructor.dept_name = dept_name  
    return d_count;  
  end
```

SQL FUNCTIONS (CONT. ,)

- Define a function that, given the name of a department, returns the count of the number of instructors in that department.
- The **function** *dept_count* can be used to find the department names and budget of all departments with more that 12 instructors.

```
select dept_name, budget  
from department  
where dept_count (dept_name ) > 12
```

SQL FUNCTIONS

- Compound statement: **begin ... end**
 - May contain multiple SQL statements between **begin** and **end**.
- **returns** -- indicates the variable-type that is returned (e.g., integer)
- **return** -- specifies the values that are to be returned as result of invoking the function
- SQL function are in fact **parameterized views** that generalize the regular notion of views by allowing parameters.

TABLE FUNCTIONS

- SQL:2003 added functions that return a relation as a result
- Example: Return all instructors in a given department

create function *instructor_of* (*dept_name* **char**(20))

returns table (

ID **varchar**(5), *name* **varchar**(20),
dept_name **varchar**(20), *salary* **numeric**(8,2))

return table

(**select** *ID*, *name*, *dept_name*, *salary*
from *instructor*

where *instructor.dept_name* = *instructor_of.dept_name*)

TABLE FUNCTIONS

Cont.,

- Usage

```
select *  
from table (instructor_of ('Music'))
```

DECLARE FUNCTION IN MYSQL

drop function if exists precourse_cnt_f;

DELIMITER \$\$

create function precourse_cnt_f(acourseid varchar(8))

returns integer

deterministic

reads sql data

begin

declare cnt integer;

select count(*) into cnt from prereq where course_id=acourseid;

return cnt;

end**\$\$**

DELIMITER ;

CALL FUNCTION IN MYSQL

- In sql functions and procedures
declare course_id varchar(10);
select '958' into course_id;
select precourse_cnt_f(course_id);

SQL PROCEDURES

- The *dept_count* function could instead be written as procedure:

```
create procedure dept_count_proc (in dept_name varchar(20),  
                                out d_count integer)
```

```
begin
```

```
    select count(*) into d_count
```

```
    from instructor
```

```
    where instructor.dept_name = dept_count_proc.dept_name
```

```
end
```


SQL PROCEDURES (CONT.,)

- The *dept_count* function could instead be written as procedure:
create or replace procedure *dept_count_proc* (in *dept_name* varchar(20),
out *d_count* integer)

Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

```
declare d_count integer;  
call dept_count_proc( 'Physics', d_count);
```

Procedures and functions can be invoked also from dynamic SQL

SQL PROCEDURES (CONT.,)

- SQL:1999 allows more than one function/procedure of the same name (called name **overloading**)
 - as long as the number of arguments differ, or at least the types of the arguments differ

LANGUAGE CONSTRUCTS FOR PROCEDURES & FUNCTIONS

- SQL supports constructs that gives it almost all the power of a general-purpose programming language.
 - Warning: most database systems implement their own variant of the standard syntax below.
- Compound statement: **begin ... end**,
 - May contain multiple SQL statements between **begin** and **end**.
 - Local variables can be declared within a compound statements

LANGUAGE CONSTRUCTS FOR PROCEDURES & FUNCTIONS (CONT.,)

- SQL supports constructs that gives it almost all the power of a general-purpose programming language.
- **While** and **repeat** statements:
 - **while** *boolean expression* **do**
sequence of statements ;
end while
 - **repeat**
sequence of statements ;
until *boolean expression*
end repeat

LANGUAGE CONSTRUCTS FOR PROCEDURES & FUNCTIONS (CONT.,)

- **For** loop
 - Permits iteration over all results of a query
- Example: Find the budget of all departments

```
declare n integer default 0;  
for r as  
    select budget from department  
do  
    set n = n + r.budget  
end for
```

LANGUAGE CONSTRUCTS FOR PROCEDURES & FUNCTIONS: IF-THEN-ELSE

- Conditional statements (**if-then-else**)
SQL:1999 also supports a **case** statement similar to C case statement
- Example procedure: registers student after ensuring classroom capacity is not exceeded
 - Returns 0 on success and -1 if capacity is exceeded
 - See book (page 177) for details

LANGUAGE CONSTRUCTS FOR PROCEDURES & FUNCTIONS : IF-THEN-ELSE

- Registers a student after ensuring classroom capacity is not exceeded
- Returns 0 on success, and -1 if capacity is exceeded.

```
create function registerStudent(  
    in s_id varchar(5),  
    in s_courseid varchar (8),  
    in s_secid varchar (8),  
    in s_semester varchar (6),  
    in s_year numeric (4,0),  
    out errorMsg varchar(100)  
  
returns integer  
begin  
    declare currEnrol int;  
    select count(*) into currEnrol  
        from takes  
        where course_id = s_courseid and sec_id = s_secid  
            and semester = s_semester and year = s_year;  
    declare limit int;  
    select capacity into limit  
        from classroom natural join section  
        where course_id = s_courseid and sec_id = s_secid  
            and semester = s_semester and year = s_year;  
    if (currEnrol < limit)  
        begin  
            insert into takes values  
                (s_id, s_courseid, s_secid, s_semester, s_year, null);  
            return(0);  
        end  
    -- Otherwise, section capacity limit already reached  
    set errorMsg = 'Enrollment limit reached for course ' || s_courseid  
        || ' section ' || s_secid;  
    return(-1);
```


LANGUAGE CONSTRUCTS FOR PROCEDURES & FUNCTIONS : EXCEPTION CONDITION 异常条件

- Signaling of exception conditions, and declaring handlers for exceptions

```
declare out_of_classroom_seats condition  
declare exit handler for out_of_classroom_seats  
begin  
...  
    .. signal out_of_classroom_seats  
end
```

- The handler here is **exit** -- causes enclosing **begin..end** to be exited
- Other actions possible on exception

LANGUAGE CONSTRUCTS FOR PROCEDURES & FUNCTIONS : EXCEPTION CONDITION 异常条件

- declaring handlers for exceptions
 - exit
 - continue
 - sqlexception
 - sqlwarning
 - not found
 - ...

EXTERNAL LANGUAGE ROUTINES

- SQL:1999 permits the use of functions and procedures written in other languages such as C or C++
- Declaring external language procedures and functions
create procedure dept_count_proc(*in dept_name* **varchar**(20),
out count **integer**)

language C

external name ' /usr/avi/bin/dept_count_proc'

create function dept_count(*dept_name* **varchar**(20))

returns integer

language C

external name ' /usr/avi/bin/dept_count'

EXTERNAL LANGUAGE ROUTINES (CONT.,)

- Benefits of external language functions/procedures:
 - more efficient for many operations, and more expressive power.

EXTERNAL LANGUAGE ROUTINES (CONT.,)

- Drawbacks

- Code to implement function may need to be loaded into database system and executed in the database system's address space.
 - risk of accidental corruption of database structures
 - security risk, allowing users access to unauthorized data
- There are alternatives, which give good security at the cost of potentially worse performance.
- Direct execution in the database system's space is used when efficiency is more important than security.

EXTERNAL LANGUAGE ROUTINES (CONT.,)

- To deal with security problems, we can do one of the following:
 - Use **sandbox** techniques
 - That is, use a safe language like Java, which cannot be used to access/damage other parts of the database code.
 - Run external language functions/procedures in a separate process, with no access to the database process' memory.
 - Parameters and results communicated via inter-process communication
- Both have performance overheads
- Many database systems support both above approaches as well as direct executing in database system address space.

MYSQL PROCEDURE

DELIMITER \$\$

```
drop procedure if exists sp_department1;  
create procedure sp_department1(out cnt integer)  
begin select count(*) into cnt from department;  
end
```

DELIMITER ;

OBJECTIVES

- Accessing SQL from a Programming Language
- Temporary Table
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP

TRIGGERS触发器

- A trigger is a statement that the system executes automatically as a side effect of a modification to the database.

TRIGGERS

- To design a trigger mechanism, we must:
 - Specify the conditions under which the trigger is to be executed.
 - Specify the actions to be taken when the trigger executes.
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.

TRIGGERING EVENTS AND ACTIONS IN SQL

- Triggering event can be **insert**, **delete** or **update**
- Triggers on update can be restricted to specific attributes
 - For example, **after update of *takes on grade***
- Values of attributes before and after an update can be referenced
 - **referencing old row as** : for deletes and updates
 - **referencing new row as** : for inserts and updates
- Triggers can be activated before an event, which can serve as extra constraints. For example, convert blank grades to null.

TRIGGERING EVENTS AND ACTIONS IN SQL

```
create trigger setnull_trigger before update of takes  
referencing new row as nrow  
for each row  
when (nrow.grade = ' ')  
    begin atomic  
        set nrow.grade = null;  
    end;
```


TRIGGER TO MAINTAIN CREDITS_EARNED

VALUE

```
• create trigger credits_earned after update of takes on (grade)
  referencing new row as nrow
  referencing old row as orow
  for each row
  when nrow.grade <> 'F' and nrow.grade is not null
    and (orow.grade = 'F' or orow.grade is null)
  begin atomic
    update student
    set tot_cred= tot_cred +
      (select credits
       from course
       where course.course_id= nrow.course_id)
    where student.id = nrow.id;
  end;
```


STATEMENT LEVEL TRIGGERS

- Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction
 - Use **for each statement** instead of **for each row**
 - Use **referencing old table** or **referencing new table** to refer to temporary tables (called **transition tables**) containing the affected rows
 - Can be more efficient when dealing with SQL statements that update a large number of rows

drop trigger if exists instuctor03_salary;

DELIMITER \$\$

create trigger instuctor03_salary before update on instructor03

for each row

begin

declare MESSAGE_TEXT varchar(200);

if NEW.salary <0 then

signal sqlstate 'AL001' set MESSAGE_TEXT = 'error salary is less than zero';

end if;

end\$\$

DELIMITER ;

DECLARE TRIGGER IN MYSQL

WHEN NOT TO USE TRIGGERS

- Triggers were used earlier for tasks such as
 - Maintaining summary data (e.g., total salary of each department)
 - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica

WHEN NOT TO USE TRIGGERS (CONT.)

- There are better ways of doing these now:
 - Databases today provide built in **materialized view** facilities to maintain summary data
 - Databases provide built-in support for replication
- Encapsulation facilities can be used instead of triggers in many cases
 - Define methods to update fields
 - Carry out actions as part of the update methods instead of through a trigger

WHEN NOT TO USE TRIGGERS (CONT.)

- **Risk of unintended execution of triggers**, for example, when
 - Loading data from a backup copy
 - Replicating updates at a remote site
 - Trigger execution can be disabled before such actions.
- Other risks with triggers:
 - Error leading to failure of critical transactions that set off the trigger
 - Cascading execution

TRIGGER

- create trigger XXX
- disable trigger XXX
- drop trigger XXX

OBJECTIVES

- Accessing SQL from a Programming Language
- Temporary Table
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP

OBJECTIVES

- Accessing SQL from a Programming Language
- Temporary Table
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP

RECURSION IN SQL

- SQL:1999 permits recursive view definition
- Example: find which courses are a prerequisite, whether directly or indirectly, for a specific course

```
create table prereq (  
    course_id      varchar(8),  
    prereq_id      varchar(8),  
    primary key (course_id, prereq_id),
```

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
BIO-399	BIO-101
CS-190	CS-101
CS-315	CS-101
CS-319	CS-101
CS-347	CS-101
EE-181	PHY-101

```
    foreign key (course_id) references course(course_id) on delete cascade,  
    foreign key (prereq_id) references course(course_id) );
```

RECURSION IN SQL:

TRANSITIVE CLOSURE传递闭包 USING ITERATION

1/3

CS-347

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
BIO-399	BIO-101
CS-190	CS-101
CS-315	CS-101
CS-319	CS-101
CS-347	CS-101
EE-181	PHY-101

Iteration Number	Tuples in c1
0	
1	(CS-301)
2	(CS-301), (CS-201)
3	(CS-301), (CS-201)
4	(CS-301), (CS-201), (CS-101)
5	(CS-301), (CS-201), (CS-101)

RECURSION IN SQL:

TRANSITIVE CLOSURE USING ITERATION 2/3

```
create function findAllPrereqs(cid varchar(8))  
    -- Finds all courses that are prerequisite (directly or indirectly) for cid  
returns table (course_id varchar(8))  
    -- The relation prereq(course_id, prereq_id) specifies which course is  
    -- directly a prerequisite for another course.  
begin  
    create temporary table c_prereq (course_id varchar(8));  
        -- table c_prereq stores the set of courses to be returned  
    create temporary table new_c_prereq (course_id varchar(8));  
        -- table new_c_prereq contains courses found in the previous iteration  
    create temporary table temp (course_id varchar(8));  
        -- table temp is used to store intermediate results  
    insert into new_c_prereq  
        select prereq_id  
        from prereq  
        where course_id = cid;
```



RECURSION IN TRANSITIVE C

```
repeat
  insert into c_prereq
    select course_id
    from new_c_prereq;

  insert into temp
    (select prereq.course_id
     from new_c_prereq, prereq
     where new_c_prereq.course_id = prereq.prereq_id
    )
  except (
    select course_id
    from c_prereq
  );
  delete from new_c_prereq;
  insert into new_c_prereq
    select *
    from temp;
  delete from temp;

until not exists (select * from new_c_prereq)
end repeat;
return table c_prereq;
end
```

RECURSION IN SQL

- **with recursive** *rec_prereq(course_id, prereq_id)* as (
 select *course_id, prereq_id*
 from *prereq*
 union
 select *rec_prereq.course_id, prereq.prereq_id,*
 from ***rec_prereq***, *prereq*
 where ***rec_prereq***.*prereq_id* = *prereq.course_id*
)
 select *
 from *rec_prereq*;

OBJECTIVES

- Accessing SQL from a Programming Language
- Temporary Table
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP

ADVANCED AGGREGATION FEATURES:

RANK排名

- Ranking is done in conjunction with an order by specification.
- Suppose we are given a relation *student_grades*(ID, GPA) giving the grade-point average of each student
- Find the rank of each student.

```
select ID, rank() over (order by GPA desc) as s_rank  
from student_grades
```

ADVANCED AGGREGATION FEATURES:

RANKING排名

- Ranking is done in conjunction with an order by specification.
- Suppose we are given a relation
student_grades(ID, GPA)
- An extra **order by** clause is needed to get them in sorted order

```
select ID, rank() over (order by GPA desc) as s_rank  
from student_grades  
order by s_rank
```

- **Ranking may leave gaps**: e.g. if 2 students have the same top GPA, both have rank 1, and the next rank is 3
- **dense_rank** does not leave gaps, so next dense rank would be 2

ADVANCED AGGREGATION FEATURES: RANKING

- Ranking can be done using basic SQL aggregation, but resultant query is very inefficient

```
select ID, (1 + (select count(*)  
                  from student_grades B  
                  where B.GPA > A.GPA)) as s_rank  
from student_grades A  
order by s_rank;
```


ADVANCED AGGREGATION FEATURES: RANKING (PARTITION)

- Ranking can be done within partition of the data.
- “Find the rank of students within each department.”

```
select ID, dept_name,  
        rank (() over (partition by dept_name order by GPA desc)  
              as dept_rank  
from dept_grades  
order by dept_name, dept_rank;
```

- Multiple **rank** clauses can occur in a single **select** clause.

ADVANCED AGGREGATION FEATURES: RANKING (PARTITION)

(Cont.,)

- Ranking is done *after* applying **group by** clause/aggregation
- Can be used to find top-n results
 - More general than the **limit** *n* clause supported by many databases, since it allows top-n within each partition

ADVANCED AGGREGATION FEATURES: RANKING

- Other ranking functions:
 - **percent_rank** (within partition, if partitioning is done)
 - **cume_dist** (cumulative distribution)
 - fraction of tuples with preceding values
 - **row_number** (non-deterministic in presence of duplicates)
- SQL:1999 permits the user to specify **nulls first** or **nulls last**

```
select ID,  
       rank ( ) over (order by GPA desc nulls last) as s_rank  
from student_grades
```


ADVANCED AGGREGATION FEATURES: RANKING

- For a given constant n , the ranking the function ***ntile(n)*** takes the tuples in each partition in the specified order, and divides them into n buckets with equal numbers of tuples.

- E.g.,

```
select ID, ntile(4) over (order by GPA desc) as quartile  
from student_grades;
```

ADVANCED AGGREGATION FEATURES:

WINDOWING 分窗

- Used to smooth out random variations.
- E.g., **moving average**: “Given sales values for each date, calculate for each date the average of the sales on that day, the previous day, and the next day” 移动平均
- **Window specification** in SQL:
 - Given relation *sales(date, value)*
select date, sum(value) over
(order by date between rows 1 preceding and 1 following)
from sales

ADVANCED AGGREGATION FEATURES: WINDOWING

- Examples of other window specifications:
 - **between rows unbounded preceding and current**
 - **rows unbounded preceding**
 - **range between 10 preceding and current row**
 - All rows with values between current row value -10 to current value
 - **range interval 10 day preceding**
 - Not including current row

ADVANCED AGGREGATION FEATURES: WINDOWING

- following

```
select year, avg(num_credits)
      over (order by year rows between 3 preceding and 2 following)
      as avg_total_credits
from tot_credits;
```

ADVANCED AGGREGATION FEATURES:

WINDOWING

- Can do windowing within partitions
- E.g., Given a relation *transaction* (*account_number*, *date_time*, *value*), where *value* is positive for a deposit and negative for a withdrawal
 - “Find total balance of each account after each transaction on the account”

```
select account_number, date_time,  
       sum (value) over  
         (partition by account_number  
          order by date_time  
          rows unbounded preceding)  
       as balance  
from transaction  
order by account_number, date_time
```

OBJECTIVES

- Accessing SQL from a Programming Language
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP

OBJECTIVES

- Accessing SQL from a Programming Language
- Temporary Table
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- **OLAP**

OLAP

- **Online Transaction Processing (OLTP)**
 - transaction
- **Online Analytical Processing (OLAP)**
 - Interactive analysis of data, allowing data to be summarized and viewed in different ways in an online fashion (with negligible delay)

EXAMPLE: SALES RELATION

<i>item_name</i>	<i>color</i>	<i>clothes_size</i>	<i>quantity</i>
skirt	dark	small	2
skirt	dark	medium	5
skirt	dark	large	1
skirt	pastel	small	11
skirt	pastel	medium	9
skirt	pastel	large	15
skirt	white	small	2
skirt	white	medium	5
skirt	white	large	3
dress	dark	small	2
dress	dark	medium	6
dress	dark	large	12
dress	pastel	small	4
dress	pastel	medium	3
dress	pastel	large	3
dress	white	small	2
dress	white	medium	3
dress	white	large	0
shirt	dark	small	2
shirt	dark	medium	6

OLAP

- Data that can be modeled as dimension attributes and measure attributes are called **multidimensional**多维 data.
 - **Measure attributes** 度量属性
 - measure some value
 - can be aggregated upon
 - e.g., the attribute *quantity* of the *sales* relation
 - **Dimension attributes** 维属性
 - define the dimensions on which measure attributes (or aggregates thereof) are viewed
 - e.g., attributes *item_name*, *color*, and *size* of the *sales* relation

CROSS TABULATION交叉表 OF SALES BY ITEM_NAME AND COLOR


clothes_size		color			
item_name	all	dark	pastel	white	total
	skirt	8	35	10	53
	dress	20	10	5	35
	shirt	14	7	28	49
	pants	20	2	5	27
	total	62	54	48	164

CROSS TABULATION交叉表 OF SALES BY ITEM_NAME AND COLOR

- The table above is an example of a **cross-tabulation (cross-tab)**交叉表, also referred to as a **pivot-table**转轴表.
 - Values for **one** of the dimension attributes form the row headers
 - Values for **another** dimension attribute form the column headers
 - **Other dimension attributes are listed on top**
 - Values in individual cells are (aggregates of) the values of the dimension attributes that specify the cell.

DATA CUBE数据立方体

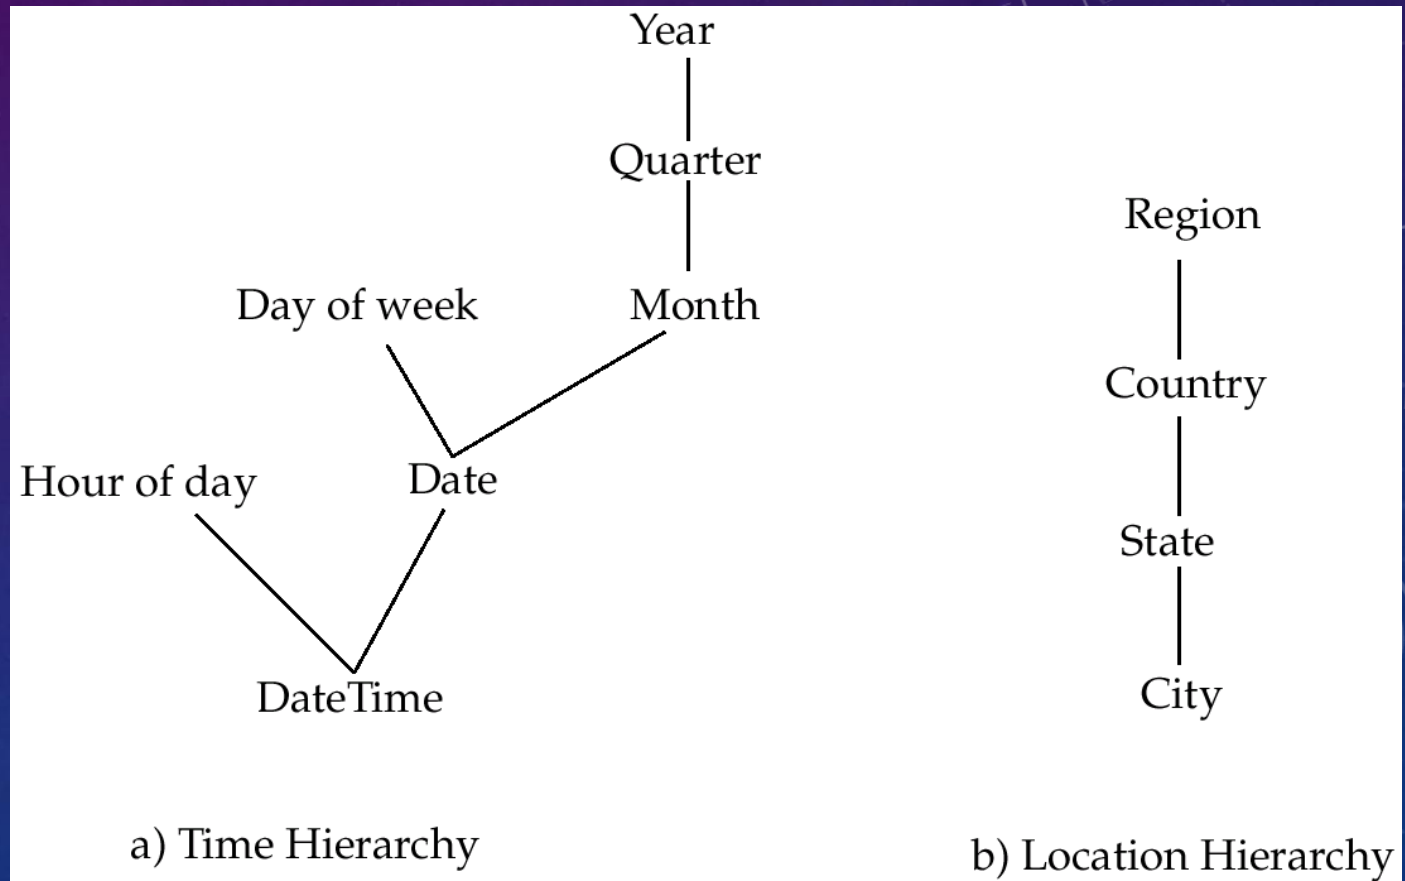
- A **data cube** is a multidimensional generalization of a **cross-tab**
- Can have ***n*** dimensions(***n***维); we show 3 below
- Cross-tabs can be used as views on a data cube



HIERARCHIES ON DIMENSIONS

- **Hierarchy** on dimension attributes: lets dimensions to be viewed at different levels of detail

E.g., the dimension DateTime can be used to aggregate by hour of day, date, day of week, month, quarter or year



CROSS TABULATION WITH HIERARCHY

- Cross-tabs can be easily extended to deal with hierarchies
 - Can drill down or roll up on a hierarchy

clothes_size: **all**

		<i>color</i>				
<i>category</i>	<i>item_name</i>	dark	pastel	white	total	
womenswear	skirt	8	8	10	53	88
	dress	20	20	5	35	
	subtotal	28	28	15		
menswear	pants	14	14	28	49	76
	shirt	20	20	5	27	
	subtotal	34	34	33		
total		62	62	48		164

RELATIONAL REPRESENTATION OF CROSS-TABS

- **Cross-tabs** can be represented as **relations**
 - We use the value **all** is used to represent aggregates.
 - The SQL standard actually uses null values in place of **all** despite confusion with regular null values.

item_name	color	clothes_size	quantity
skirt	dark	all	8
skirt	pastel	all	35
skirt	white	all	10
skirt	all	all	53
dress	dark	all	20
dress	pastel	all	10
dress	white	all	5
dress	all	all	35
shirt	dark	all	14
shirt	pastel	all	7
shirt	White	all	28
shirt	all	all	49
pant	dark	all	20
pant	pastel	all	2
pant	white	all	5
pant	all	all	27
all	dark	all	62
all	pastel	all	54
all	white	all	48
all	all	all	164

PIVOT转轴

```
select *  
from sales  
pivot (  
    sum(quantity)  
    for color in ('dark','pastel','white')  
)  
order by item_name;
```

item_name	clothes_size	dark	pastel	white
skirt	small	2	11	2
skirt	medium	5	9	5
skirt	large	1	15	3
dress	small	2	4	2
dress	medium	6	3	3
dress	large	12	3	0
shirt	small	2	4	17
shirt	medium	6	1	1
shirt	large	6	2	10
pants	small	14	1	3
pants	medium	6	0	0
pants	large	0	1	2

EXTENDED AGGREGATION TO SUPPORT OLAP

- The **cube** operation computes union of **group by**'s on every subset of the specified attributes
- Example relation for this section:

sales(item_name, color, clothes_size, quantity)

```
select item_name, sum(quantity)
from sales
group by item_name;
```

item_name	quantity
skirt	53
dress	35
shirt	49
pants	27

EXTENDED AGGREGATION TO SUPPORT OLAP (CONT. ,)

- The **cube** operation computes union of **group by**'s on every subset of the specified attributes

- Example relation for this section:

sales(item_name, color, clothes_size, quantity)

- E.g. consider the query

```
select item_name, color, clothes_size, sum(quantity)  
from sales  
group by cube(item_name, color, clothes_size)
```

EXTENDED AGGREGATION TO (CONT. ,)

```
select item_name, color,  
clothes_size, sum(quantity)  
from sales  
group by cube(item_name, color,  
clothes_size)
```

<i>item_name</i>	<i>color</i>	<i>clothes_size</i>	<i>quantity</i>
skirt	dark	all	8
skirt	pastel	all	35
skirt	white	all	10
skirt	all	all	53
dress	dark	all	20
dress	pastel	all	10
dress	white	all	5
dress	all	all	35
shirt	dark	all	14
shirt	pastel	all	7
shirt	White	all	28
shirt	all	all	49
pant	dark	all	20
pant	pastel	all	2
pant	white	all	5
pant	all	all	27
all	dark	all	62
all	pastel	all	54
all	white	all	48
all	all	all	164

EXTENDED AGGREGATION TO SUPPORT OLAP (CONT. ,)

- This computes the union of eight different groupings of the *sales* relation:

$\{ (item_name, color, size), (item_name, color),$
 $(item_name, size), (color, size),$
 $(item_name), (color),$
 $(size), () \}$

where () denotes an empty **group by** list.

- For each grouping, the result contains the null value for attributes not present in the grouping.

ONLINE ANALYTICAL PROCESSING OPERATIONS: GROUPING

- Relational representation of cross-tab that we saw earlier, but with *null* in place of **all**, can be computed by

```
select item_name, color, sum(number)  
from sales  
group by cube(item_name, color)
```

ONLINE ANALYTICAL PROCESSING OPERATIONS: GROUPING (CONT.,)

- The function **grouping()** can be applied on an attribute
 - Returns 1 if the value is a null value representing all, and returns 0 in all other cases.

```
select item_name, color, size, sum(number),  
       grouping(item_name) as item_name_flag,  
       grouping(color) as color_flag,  
       grouping(size) as size_flag,  
from sales  
group by cube(item_name, color, size)
```

ONLINE ANALYTICAL PROCESSING OPERATIONS: GROUPING (CONT.,)

- Can use the function **decode()** in the **select** clause to replace such nulls by a value such as **all**
 - `decode (value, match-1, replacement-1, match-2, replacement-2, ..., match-N, replacement-N, default-replacement);`
 - E.g., replace *item_name* in first query by **decode(grouping(*item_name*), 1, 'all', *item_name*)**

ONLINE ANALYTICAL PROCESSING OPERATIONS: ROLLUP 上卷

- The **rollup** construct generates union on every prefix of specified list of attributes 由细粒度到粗粒度

- E.g.,

```
select item_name, color, size, sum(number)  
from sales  
group by rollup(item_name, color, size)
```

Generates union of four groupings:

```
{ (item_name, color, size), (item_name, color), (item_name), ( ) }
```

ONLINE ANALYTICAL PROCESSING OPERATIONS: ROLLUP (CONT. ,)

- Rollup can be used to generate aggregates at multiple levels of a hierarchy.
- E.g., suppose table *itemcategory*(*item_name*, *category*) gives the category of each item. Then

```
select category, item_name, sum(number)  
from sales, itemcategory  
where sales.item_name = itemcategory.item_name  
group by rollup(category, item_name)
```

would give a hierarchical summary by *item_name* and by *category*.

ONLINE ANALYTICAL PROCESSING OPERATIONS

- Multiple rollups and cubes can be used in a single group by clause
 - Each generates set of group by lists, cross product of sets gives overall set of group by lists
- E.g.,

```
select item_name, color, size, sum(number)  
from sales  
group by rollup(item_name), rollup(color, size)
```

generates the groupings

$\{item_name, ()\} \times \{(color, size), (color), ()\}$

$= \{ (item_name, color, size), (item_name, color), (item_name), (color, size), (color), () \}$

ONLINE ANALYTICAL PROCESSING OPERATIONS

- **Pivoting** 转轴: changing the dimensions used in a cross-tab is called
- **Slicing** 分片: creating a cross-tab for fixed values only
 - Sometimes called **dicing** 切块, particularly when values for multiple dimensions are fixed.
- **Rollup** 上卷: moving from finer-granularity data to a coarser granularity
- **Drill down** 下钻: The opposite operation - that of moving from coarser-granularity data to finer-granularity data

OLAP IMPLEMENTATION

- The earliest OLAP systems used multidimensional arrays in memory to store data cubes, and are referred to as **multidimensional OLAP (MOLAP)** systems.
- OLAP implementations using only relational database features are called **relational OLAP (ROLAP)** systems
- Hybrid systems, which store some summaries in memory and store the base data and other summaries in a relational database, are called **hybrid OLAP (HOLAP)** systems.

- Early OLAP systems precomputed *all* possible aggregates in order to provide online response
 - Space and time requirements for doing so can be very high
 - 2^n combinations of **group by**
 - It suffices to precompute some aggregates, and compute others on demand from one of the precomputed aggregates
 - Can compute aggregate on (*item_name, color*) from an aggregate on (*item_name, color, size*)
 - For all but a few “non-decomposable” aggregates such as median
 - is cheaper than computing it from scratch

- Several optimizations available for computing multiple aggregates
 - Can compute aggregate on (*item_name*, *color*) from an aggregate on (*item_name*, *color*, *size*)
 - Can compute aggregates on (*item_name*, *color*, *size*), (*item_name*, *color*) and (*item_name*) using a single sorting of the base data

SUMMARY

- Accessing SQL from a Programming Language
- Temporary Table
- Functions and Procedures
- Triggers
- Recursive Queries
- Advanced Aggregation Features
- OLAP



Q&A?



THANKS !

leexudong@nankai.edu.cn